

OPENTSDDB

2.4 and 3.0 Update

Who Am I?

Chris Larsen

- Maintainer and author for OpenTSDB since 2013
- Software Engineer @ Yahoo
- Central Monitoring Team

Who I'm not:

- A marketer
- A sales person

What Is OpenTSDB?

- Open Source Time Series Database
- Scales to 10s of millions of writes per second
- Sucks up *all* data and keeps going
- Never lose precision (if you have space)
- Scales using HBase or Bigtable



What are Time Series?



- **Time Series:** A sequence of discrete data points (values) ordered and indexed by time associated with an identity.

E.g.:

web01.sys.cpu.busy.pct 45% 1/1/207 12:01:00

web01.sys.cpu.busy.pct 52% 1/1/207 12:02:00


web01.sys.cpu.busy.pct 35% 1/1/207 12:03:00

^ Identity

^ Value ^ Timestamp

What are Time Series?


```
1 [{  
2   "metric": "system.cpu.busy.pct",  
3   "tags": {  
4     "_aggregate": "raw",  
5     "host": "web01"  
6   },  
7   "dps": {  
8     "1486619640": 5.2,  
9     "1486619700": 5.5,  
10    "1486619760": 6.1,  
11    "1486619820": 7.3,  
12    "1486619880": 7.6  
13  }  
14 }]
```



What are Time Series?

Data Point:

- Metric + Tags
- + Value: 42
- + Timestamp: 1234567890



^ a data point ^

sys.cpu.user 1234567890 42 host=web01 cpu=0

- Payload could also be a string, a blob, a histogram, etc.

What are HBase and Bigtable?

- HBase is an OSS distributed LSM backed hash table based on Google's Bigtable.
- Key value, row based column store.
- Sorted by row, columns and cell versions.
- Supports:
 - Scans across rows with filters.
 - Get specific row and/or columns.
 - Atomic operations.
- CP from CAP theorem.



APACHE
HBASE



OpenTSDB Schema

- Row key is a concatenation of UUIDs and time:
 - salt + metric + timestamp + tagk1 + tagv1... + tagkN + tagvN
- `sys.cpu.user 1234567890 42 host=web01 cpu=0`
`\x01\x00\x00\x01\x49\x95\xFB\x70\x00\x00\x01\x00\x00\x01\x00\x00\x02\x00\x00\x02`
- Timestamp normalized on hour or daily boundaries.
- All data points for an hour or day are stored in one row.
- Data: VLE 64 bit signed integers or single/double precision signed floats, Strings and raw histograms.
- Saves storage space but requires UUID conversion.

OpenTSDB Schema

Row Key	Columns (qualifier/value)
m t1 tagk1 tagv1	o1/v1 o2/v2 o3/v3
m t1 tagk1 tagv2	o1/v1 o2/v2
m t1 tagk1 tagv1 tagk2 tagv3	o1/v1 o2/v2 o3/v3
m t1 tagk1 tagv2 tagk2 tagv4	o1/v1 o3/v3
m t1 tagk3 tagv5	o1/v1 o2/v2 o3/v3
m t1 tagk3 tagv6	o2/v2
m t2 tagk1 tagv1	o1/v1 o3/v3
m t2 tagk1 tagv2	o1/v1 o2/v2

```
1 {  
2     "start": "t1",  
3     "end": "t2",  
4     "queries": [{  
5         "metric": "m1",  
6         "tags": {  
7             "tagk1": "tagv1"  
8         },  
9         "explicitTags": true  
10    }  
11  
12 ]  
13 }
```

OpenTSDB Use Cases

- Backing store for Argus: Open source monitoring and alerting system.
- 50M writes per minute.
- ~4M writes per TSD per minute.
- 23k queries per minute.
- <https://github.com/salesforce/Argus>



OpenTSDB Use Cases

The Yahoo! logo is displayed in a stylized, purple, serif font. It is positioned in the upper right corner of the slide, enclosed within a light gray rectangular box.

- Monitoring system, network and application performance and statistics.
- Single cluster: 10M to 18M writes/s ~ 3PB.
- Multi-tenant and Kerberos secure HBase.
- ~200k writes per second per TSD.
- Central monitoring for all Yahoo properties.
- Over 1 billion active time series served.
- Leading committer to OpenTSDB.


Other Users

EMC²

StackExchange 



Bloomberg



YAHOO!

New for OpenTSDB 2.4

- Rollup / Pre-Aggregated storage and querying
 - Improves query speed
 - Allows for high-resolution data to be TTL'd out
- Histogram/Digests/Sketches
 - Accurate percentile calculations on distributed measurements such as latencies.
- Date Tiered Compaction support
- Authentication/Authorization plugin

The Problem of Percentiles

- Aggregating percentiles ==

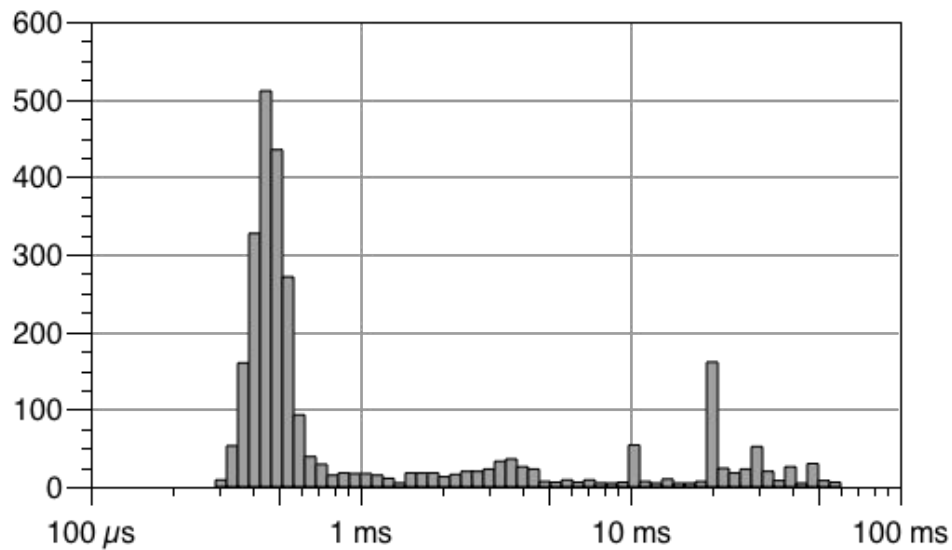
```
latency.p99.9 42.50 host=web01  
latency.p99.9 58.98 host=web02  
latency.p99.9 41.28 host=web03  
latency.p99.9 41.94 host=web04
```

- Averaging percentiles is inaccurate.
E.g. **46.175** hides the bad host, **web02**
- Max is more useful for finding bad hosts
- But there are better ways...



Histograms

- Distribution of frequency of measurements over a time period
- Simplest form: latency measurement buckets storing counts falling within those buckets. E.g.



```
latency.histogram 0,15.0=0:15.0,30.0=1:30.0,45.0=4:45.0,60.0=0 host=web01
latency.histogram 0,15.0=1:15.0,30.0=0:30.0,45.0=2:45.0,60.0=4 host=web02
latency.histogram 0,15.0=2:15.0,30.0=0:30.0,45.0=4:45.4,60.0=0 host=web03
latency.histogram 0,15.0=0:15.0,30.0=1:30.0,45.0=4:45.0,60.0=0 host=web04
```

Histograms

Histogram	p99	p85	p50
latency.histogram 0,15.0= 0 :15.0,30.0= 1 :30.0,45.0= 4 :45.0,60.0= 0 host=web01	37.5	37.5	37.5
latency.histogram 0,15.0= 1 :15.0,30.0= 0 :30.0,45.0= 2 :45.0,60.0= 4 host=web02	52.5	52.5	52.5
latency.histogram 0,15.0= 2 :15.0,30.0= 0 :30.0,45.0= 4 :45.4,60.0= 0 host=web03	37.5	37.5	37.5
latency.histogram 0,15.0= 0 :15.0,30.0= 1 :30.0,45.0= 4 :45.0,60.0= 0 host=web04	37.5	37.5	37.5
Averaged Percentiles:	41.25	41.25	41.25
Summed Histograms:			
latency.histogram 0,15.0= 3 :15.0,30.0= 2 :30.0,45.0= 14 :45.0,60.0= 4	52.5	52.5	37.5

Histograms

- Pros:
 - Fixed size (877 bytes for 97 buckets per data point)
 - Richer analysis (probability distribution, etc)
 - Mergable via group by and downsampling
 - Fixed rank error, variable value error
- Cons:
 - Much more network/storage space required
 - Loss of accuracy (somewhere within the bucket) but precise
 - Common metrics libraries lack support

Pluggable Implementations

Yahoo's Data Sketches

- Collection of approximation algorithms with mergability and configurable accuracy v. size (~26k for 2M measurements)
- Deterministic rank error
- Tapering log size with N measurements per sketch
- Good for median percentiles
- <https://datasketches.github.io/>



Pluggable Implementations


T-Digest

- Offshoot of Q-Digest K-means clustering quantile approximations
- Small error at top and bottom of the quantile range
- Mergable
- Able to store floating point as well as integers
- <https://github.com/tdunning/t-digest>

The Problem of Appends

- 2.2 Introduced appends to move away from TSD compactions.
- 1 second resolution = 3600 columns per row => compact into 1.
- But with appends, HBase:
 - Reads the column (from memstore or disk)
 - Appends the data and writes back to memstore (and possibly block cache)
 - Send full data back to the client

What is
our appendix
for?



The Problem of Appends

- Negatives:
 - Possible disk thrashing if columns have been compacted out of the memstore
 - Higher CPU utilization on the region servers
 - Longer wait time on the client side
- Future Solution:
 - Yahoo's HBase developers (Francis, Thiruvell) working on an optimization using coprocessors.
 - Trials underway, details in August

The Problem of Compactions

- HBase compaction merges multiple store files into one, saving space.
- But if we assume the data is time series, with older data immutable and updates only to new data...
- ...we can avoid re-compacting old files that won't change and skip them at scan time.
- HBASE-15181 from Yahoo and Flurry supports organizing store files by date and time.
- PR #990 from Karan at Salesforce allows TSDB to write HBase timestamps

AsyncHBase 1.8

- AsyncHBase is a fully asynchronous, multi-threaded HBase client
- Supports HBase 0.90 to 1.x
- Faster and less resource intensive than the native HBase client
- Support for scanner filters, META prefetch, “fail-fast” RPCs



AsyncHBase 1.8

- Batched GetRequests thanks to Tian-Ying at Pinterest and Bizu at Yahoo
- Reverse scanning support thanks to Jiayun at Harvard
- HBase 1.3.x+ support thanks to Karan at SalesForce
- MultipleColumnPrefixFilter
- Skip WAL with increments
- AtomicIncrements with multiple columns per request

OpenTSDB on Bigtable



- Bigtable
 - Hosted Google Service
 - Client uses HTTP2 and GRPC for communication
- OpenTSDB heads home
 - Based on a time series store on Bigtable at Google
 - Identical schema as HBase
 - Same filter support (fuzzy filters are coming)



Google Cloud Platform

OpenTSDB 3.0



- **Problem:** Queries are slow and the order of operations is immutable
- **Solutions:** *(This part is ready for testing!)*
 - New composable query layer allowing operations in any order
 - Support for querying multiple sources and merging the results (e.g. use Facebook's Berengi as a write-cache and Redis as a query cache)
 - Support for multi-cluster queries for active-active, high-availability setups

OpenTSDB 3.0


- **Problem:** Storing other types of data or using other backends is a pain.
- **Solutions:** *(In progress)*
 - Pluggable storage interface allowing for various schemas and implementations (e.g. native HTable client, AsyncHBase, native Bigtable client, etc)
 - Abstracted data types for pluggable implementations of time series (e.g. raw binary, histograms, SCADA data)

OpenTSDB 3.0

- **Problem:** What about anomaly detection, forecasting, etc?
- **Solutions:** *(In progress)*
 - Integration with Yahoo's EGADS time series functions library
 - Period-over-period analysis with model caching
 - Clustering algorithms for detecting outliers
 - <https://github.com/yahoo/egads>

OpenTSDB 3.0

- New Java APIs
- Servlet for standard deployment using your favorite server
- Tracing with Zipkin and OpenTracing
- New debugging UI
- Improved Docker support



Alternative TSDBs



DalmatinerDB



TIMESCALE



AXIBASE



Warp 10



BTrDB: Berkeley Tree Database

<https://misfra.me/2016/04/09/tsdb-list/>

More Info and Credits

- Thanks to the Monitoring and HBase teams at Yahoo, Pythian for Bigtable support and our OSS contributors!
- Contribute at github.com/OpenTSDB/opentsdb
- Website: opentsdb.net
- Mailing List: groups.google.com/group/opentsdb

Images

- https://commons.wikimedia.org/wiki/File:Programmer_writing_code_with_Unit_Tests.jpg
- <http://www.doncio.navy.mil/CHIPS/ArticleDetails.aspx?ID=8098>
- <https://i.imgur.com/t96s8.jpg>
- https://commons.wikimedia.org/wiki/File:Twemoji_1f626.svg
- <https://xkcd.com/1425/>
- https://commons.wikimedia.org/wiki/Emoji#/media/File:Twemoji_1f623.svg
- https://c1.staticflickr.com/1/508/32307332875_40e73bf750_b.jpg
- <http://code.flickr.net/2008/10/27/counting-timing/>
- <http://3.bp.blogspot.com/-tTXEI5liQh4/VQqaJz4LtSI/AAAAAAAAEL8/n5AwTVNI-U/s1600/Introduction%2Bto%2BSQL.png>